

Label-efficient Land Cover Classification at High Spatial Resolution using Multi-stage Pre-training

Dakota Hester

Mississippi State University
Department of Agricultural and Biological Engineering

February 3, 2025

Land Use and Land Cover (LULC)

- **Land cover:** biophysical material covering the Earth's surface (e.g., water, vegetation, soil, structures, etc.) [1]
"What is there?"




Figure 1: Land cover map of Mississippi.

Land Use and Land Cover (LULC)

- ▶ **Land cover:** biophysical material covering the Earth's surface (e.g., water, vegetation, soil, structures, etc.) [1]
“What is there?”
- ▶ **Land use:** human interaction with the land cover (e.g., agriculture, forest practices, urban development, etc.) [2, 3]
“How is it being used?”




Figure 1: Land cover map of Mississippi.

Land Use and Land Cover (LULC)

- ▶ **Land cover:** biophysical material covering the Earth's surface (e.g., water, vegetation, soil, structures, etc.) [1]
“What is there?”
- ▶ **Land use:** human interaction with the land cover (e.g., agriculture, forest practices, urban development, etc.) [2, 3]
“How is it being used?”
- ▶ LULC data are used in a wide range of disciplines, including agriculture, forestry, urban planning, and environmental science [4–6]




Figure 1: Land cover map of Mississippi.

Pixel-wise Classification

- ▶ Traditional classifiers usually operate by transforming a feature vector into output class probabilities for a given pixel.

Pixel-wise Classification

- ▶ Traditional classifiers usually operate by transforming a feature vector into output class probabilities for a given pixel.
- ▶ Common data sources include:

Pixel-wise Classification

- ▶ Traditional classifiers usually operate by transforming a feature vector into output class probabilities for a given pixel.
- ▶ Common data sources include:
 - ▶ Surface reflectance data

Pixel-wise Classification

- ▶ Traditional classifiers usually operate by transforming a feature vector into output class probabilities for a given pixel.
- ▶ Common data sources include:
 - ▶ Surface reflectance data
 - ▶ Derived spectral indices (i.e., NDVI, EVI, NDWI, etc.)

Pixel-wise Classification

- ▶ Traditional classifiers usually operate by transforming a feature vector into output class probabilities for a given pixel.
- ▶ Common data sources include:
 - ▶ Surface reflectance data
 - ▶ Derived spectral indices (i.e., NDVI, EVI, NDWI, etc.)
 - ▶ Ancillary information (i.e., topography, soil, etc) [6, 7]

Pixel-wise Classification

- ▶ Traditional classifiers usually operate by transforming a feature vector into output class probabilities for a given pixel.
- ▶ Common data sources include:
 - ▶ Surface reflectance data
 - ▶ Derived spectral indices (i.e., NDVI, EVI, NDWI, etc.)
 - ▶ Ancillary information (i.e., topography, soil, etc) [6, 7]
- ▶ Non-parametric, supervised classifiers are often used (e.g., Decision Trees, Random Forests, Support Vector Machines, etc.) [8, 9].

Pixel-wise Classification




Figure 2: Pixel-wise classification workflow.

Land Cover at High Spatial Resolution




Figure 3: Comparison of land cover maps at 30 m (left) and 1 m (right) spatial resolution.

Decreased Class Discriminability

Classifying pixels gets **more difficult** as spatial resolution **increases** [10, 11].




Figure 4: As spatial resolution increases, inter-class variance decreases [11].

Decreased Class Discriminability




Figure 5: Applying a pixel-wise classifier to high spatial resolution data typically results in noisy classification maps.

“Why do we need deep models?”

- ▶ Deep learning models have several key advantages over traditional classifiers:

“Why do we need deep models?”

- ▶ Deep learning models have several key advantages over traditional classifiers:
 1. **Feature extraction** — deep models can learn complex hierarchies of **spatial** features that are difficult to engineer by hand [12].

“Why do we need deep models?”

- ▶ Deep learning models have several key advantages over traditional classifiers:
 1. **Feature extraction** — deep models can learn complex hierarchies of **spatial** features that are difficult to engineer by hand [12].
 2. **Flexibility** — the nature of deep learning architectures allows them to be applied to a wide range of tasks across different modalities, and can be modified to best suit a given task [13].

“Why do we need deep models?”

- ▶ Deep learning models have several key advantages over traditional classifiers:
 1. **Feature extraction** — deep models can learn complex hierarchies of **spatial** features that are difficult to engineer by hand [12].
 2. **Flexibility** — the nature of deep learning architectures allows them to be applied to a wide range of tasks across different modalities, and can be modified to best suit a given task [13].
 3. **End-to-end learning** — unlike traditional methods, deep models can be trained in an end-to-end fashion, where little to no pre-processing or post-processing is required [14–17].

Convolutional Neural Networks (CNNs)




Figure 6: A ResNet-18 CNN architecture [18]. A sampling of learned convolutional filters from the first layer is shown.

U-Net




Figure 7: U-Net architecture [19].

Data Requirements

- ▶ While deep learning networks are powerful, their primary drawback is the amount of data required to train them.

Data Requirements

- ▶ While deep learning networks are powerful, their primary drawback is the amount of data required to train them.
- ▶ For example, ResNet-50, a popular CNN architecture, has over 25,000,000 parameters that need to be learned [18].

Data Requirements

- ▶ While deep learning networks are powerful, their primary drawback is the amount of data required to train them.
- ▶ For example, ResNet-50, a popular CNN architecture, has over 25,000,000 parameters that need to be learned [18].
- ▶ For land cover, this is especially problematic, as high resolution data is complex and difficult to annotate at scale [20].

Data Requirements

- ▶ While deep learning networks are powerful, their primary drawback is the amount of data required to train them.
- ▶ For example, ResNet-50, a popular CNN architecture, has over 25,000,000 parameters that need to be learned [18].
- ▶ For land cover, this is especially problematic, as high resolution data is complex and difficult to annotate at scale [20].
- ▶ One approach to mitigate this issue is to use **transfer learning**, where a network can be **pre-trained** on a related task with a large amount of data, then **fine-tuned** on the target task with a smaller amount of data [21].

Transfer Learning




Figure 8: Transfer learning paradigm.

Denoising Autoencoder

- ▶ While labeled data is optimal for pre-training, some pre-training tasks can be performed in a **self-supervised** manner, where the network learns to predict some property of the input data without requiring labels [22].

Denoising Autoencoder

- ▶ While labeled data is optimal for pre-training, some pre-training tasks can be performed in a **self-supervised** manner, where the network learns to predict some property of the input data without requiring labels [22].
- ▶ Denoising autoencoders are a type of CNN that are trained to remove noise from an input image [23].

Denoising Autoencoder

- ▶ While labeled data is optimal for pre-training, some pre-training tasks can be performed in a **self-supervised** manner, where the network learns to predict some property of the input data without requiring labels [22].
- ▶ Denoising autoencoders are a type of CNN that are trained to remove noise from an input image [23].
- ▶ In effect, noise is added to an input image, then train the network to predict the original image from the noisy image.

Denoising Autoencoder

- ▶ While labeled data is optimal for pre-training, some pre-training tasks can be performed in a **self-supervised** manner, where the network learns to predict some property of the input data without requiring labels [22].
- ▶ Denoising autoencoders are a type of CNN that are trained to remove noise from an input image [23].
- ▶ In effect, noise is added to an input image, then train the network to predict the original image from the noisy image.
- ▶ The network needs to learn **features** that are **robust** to noise while being meaningful for reconstruction [24].

Our Goal

- ▶ Our overarching goal is to develop a 1m resolution land cover map that covers the entire state of Mississippi with a high degree of accuracy.

Our Goal

- ▶ Our overarching goal is to develop a 1m resolution land cover map that covers the entire state of Mississippi with a high degree of accuracy.
- ▶ Data annotation is expensive and time-consuming, so the main focus of our work is to devise a creative pre-training strategy that combines self-supervised pre-training with a supervised fine-tuning stage to reduce the amount of labeled data required for training a deep model.

Our Goal

- ▶ Our overarching goal is to develop a 1m resolution land cover map that covers the entire state of Mississippi with a high degree of accuracy.
- ▶ Data annotation is expensive and time-consuming, so the main focus of our work is to devise a creative pre-training strategy that combines self-supervised pre-training with a supervised fine-tuning stage to reduce the amount of labeled data required for training a deep model.
- ▶ We construct a large U-Net model that undergoes various pre-training regimens, then fine-tune the model on only 119 tiles of fully annotated data from the state of Mississippi.

Our Goal

- ▶ Our overarching goal is to develop a 1m resolution land cover map that covers the entire state of Mississippi with a high degree of accuracy.
- ▶ Data annotation is expensive and time-consuming, so the main focus of our work is to devise a creative pre-training strategy that combines self-supervised pre-training with a supervised fine-tuning stage to reduce the amount of labeled data required for training a deep model.
- ▶ We construct a large U-Net model that undergoes various pre-training regimens, then fine-tune the model on only 119 tiles of fully annotated data from the state of Mississippi.
- ▶ Transfer learning is not an uncommon approach, but self-supervised pre-training is not common in the remote sensing community — especially in the context of land cover classification. As such, we are most interested in the performance of our self-supervised model.

Training Strategy




Figure 9: Model training flowchart.

Pre-training Strategies

Pre-training with a denoising autoencoder improved the performance of the model under the 2-stage pre-training strategy.




Figure 10: Performance of models on the Mississippi land cover dataset.

Pre-training Strategies

An additional fully supervised pre-training stage substantially improves the performance of the model.




Figure 11: Performance of the best model with and without additional supervised pre-training.

Confusion Matrix




Figure 12: Confusion matrix of the final model on the Mississippi land cover dataset.

Warren County Land Cover



Figure 13: Warren County Land Cover Map.

Summary

- ▶ Large scale pre-training of a deep model using an unsupervised regimen can enable the training of an over-parameterized deep model on a small amount of labeled data.

Summary

- ▶ Large scale pre-training of a deep model using an unsupervised regimen can enable the training of an over-parameterized deep model on a small amount of labeled data.
- ▶ Further, additional pre-training using a modest amount of labeled data can boost the performance of the model even further.

Summary

- ▶ Large scale pre-training of a deep model using an unsupervised regimen can enable the training of an over-parameterized deep model on a small amount of labeled data.
- ▶ Further, additional pre-training using a modest amount of labeled data can boost the performance of the model even further.
- ▶ Still, classifying more visually intricate land cover classes remains a challenge, even with a large model.

Future Works

- ▶ More data — clear need to improve the accuracy of cultivated crops and barren land classes.

Future Works

- ▶ More data — clear need to improve the accuracy of cultivated crops and barren land classes.
- ▶ Post-processing using OBIA-inspired techniques to refine classification map edges.

Future Works

- ▶ More data — clear need to improve the accuracy of cultivated crops and barren land classes.
- ▶ Post-processing using OBIA-inspired techniques to refine classification map edges.
- ▶ Applying model to super-resolved satellite imagery.

Acknowledgements



References I

- [1] John R Townshend. "Land Cover". In: *International Journal of Remote Sensing* 13.6-7 (Apr. 1992), pp. 1319–1328. ISSN: 0143-1161. DOI: 10.1080/01431169208904193. (Visited on 08/20/2024).
- [2] James R Anderson. "Land-Use Classification Schemes". In: *Photogrammetric Engineering* (1971).
- [3] James Richard Anderson. *A Land Use and Land Cover Classification System for Use with Remote Sensor Data*. Vol. 964. US Government Printing Office, 1976.
- [4] T. R. Loveland et al. "Development of a Global Land Cover Characteristics Database and IGBP DISCover from 1 Km AVHRR Data". In: *International Journal of Remote Sensing* 21.6-7 (Jan. 2000), pp. 1303–1330. ISSN: 0143-1161. DOI: 10.1080/014311600210191. (Visited on 09/05/2024).

References II

- [5] George Grekousis, Giorgos Mountrakis, and Marinos Kavouras. “An Overview of 21 Global and 43 Regional Land-Cover Mapping Products”. In: *International Journal of Remote Sensing* 36.21 (Nov. 2015), pp. 5309–5335. ISSN: 0143-1161. DOI: 10.1080/01431161.2015.1093195. (Visited on 09/05/2024).
- [6] Michael A. Wulder et al. “Land Cover 2.0”. In: *International Journal of Remote Sensing* 39.12 (June 2018), pp. 4254–4284. ISSN: 0143-1161, 1366-5901. DOI: 10.1080/01431161.2018.1452075. (Visited on 07/03/2023).
- [7] Taskin Kavzoglu and Furkan Bilucan. “Effects of Auxiliary and Ancillary Data on LULC Classification in a Heterogeneous Environment Using Optimized Random Forest Algorithm”. In: *Earth Science Informatics* 16.1 (Mar. 2023), pp. 415–435. ISSN: 1865-0481. DOI: 10.1007/s12145-022-00874-9. (Visited on 09/09/2024).

References III

- [8] Pall Oskar Gislason, Jon Atli Benediktsson, and Johannes R. Sveinsson. “Random Forests for Land Cover Classification”. In: *Pattern Recognition Letters*. Pattern Recognition in Remote Sensing (PRRS 2004) 27.4 (Mar. 2006), pp. 294–300. ISSN: 0167-8655. DOI: 10.1016/j.patrec.2005.08.011. (Visited on 09/12/2024).
- [9] J. R. BAKER et al. “Advances in Classification for Land Cover Mapping Using SPOT HRV Imagery”. In: *International Journal of Remote Sensing* 12.5 (May 1991), pp. 1071–1085. ISSN: 0143-1161. DOI: 10.1080/01431169108929711. (Visited on 09/12/2024).
- [10] B. L. Markham and J. R. G. Townshend. *Land Cover Classification Accuracy as a Function of Sensor Spatial Resolution*. Jan. 1981. (Visited on 09/09/2024).

References IV

- [11] Curtis E. Woodcock and Alan H. Strahler. “The Factor of Scale in Remote Sensing”. In: *Remote Sensing of Environment* 21.3 (Apr. 1987), pp. 311–332. ISSN: 00344257. DOI: 10.1016/0034-4257(87)90015-0. (Visited on 07/31/2024).
- [12] Manjunath Jogin et al. “Feature Extraction Using Convolution Neural Networks (CNN) and Deep Learning”. In: *2018 3rd IEEE International Conference on Recent Trends in Electronics, Information & Communication Technology (RTEICT)*. May 2018, pp. 2319–2323. DOI: 10.1109/RTEICT42901.2018.9012507. (Visited on 01/20/2025).
- [13] Danfeng Hong et al. “More Diverse Means Better: Multimodal Deep Learning Meets Remote-Sensing Imagery Classification”. In: *IEEE Transactions on Geoscience and Remote Sensing* 59.5 (May 2021), pp. 4340–4354. ISSN: 1558-0644. DOI: 10.1109/TGRS.2020.3016820. (Visited on 01/20/2025).

References V

- [14] Mohammad D. Hossain and Dongmei Chen. “Segmentation for Object-Based Image Analysis (OBIA): A Review of Algorithms and Challenges from Remote Sensing Perspective”. In: *ISPRS Journal of Photogrammetry and Remote Sensing* 150 (Apr. 2019), pp. 115–134. ISSN: 0924-2716. DOI: [10.1016/j.isprsjprs.2019.02.009](https://doi.org/10.1016/j.isprsjprs.2019.02.009). (Visited on 08/01/2024).
- [15] I. Kotaridis and M. Lazaridou. “Remote Sensing Image Segmentation Advances: A Meta-Analysis”. In: *ISPRS Journal of Photogrammetry and Remote Sensing* 173 (2021), pp. 309–322. DOI: [10.1016/j.isprsjprs.2021.01.020](https://doi.org/10.1016/j.isprsjprs.2021.01.020).
- [16] Bipul Neupane, Teerayut Horanont, and Jagannath Aryal. “Deep Learning-Based Semantic Segmentation of Urban Features in Satellite Images: A Review and Meta-Analysis”. In: *Remote Sensing* 13.4 (Jan. 2021), p. 808. ISSN: 2072-4292. DOI: [10.3390/rs13040808](https://doi.org/10.3390/rs13040808). (Visited on 10/07/2024).

References VI

- [17] Ning Zang et al. “Land-Use Mapping for High-Spatial Resolution Remote Sensing Image Via Deep Learning: A Review”. In: *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing* 14 (2021), pp. 5372–5391. ISSN: 2151-1535. DOI: 10.1109/JSTARS.2021.3078631. (Visited on 10/07/2024).
- [18] Kaiming He et al. *Deep Residual Learning for Image Recognition*. Dec. 2015. arXiv: 1512.03385 [cs]. (Visited on 07/17/2024).
- [19] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. *U-Net: Convolutional Networks for Biomedical Image Segmentation*. 2015. DOI: 10.48550/ARXIV.1505.04597. (Visited on 07/17/2024).
- [20] Yuchi Ma et al. “Transfer Learning in Environmental Remote Sensing”. In: *Remote Sensing of Environment* 301 (Feb. 2024), p. 113924. ISSN: 0034-4257. DOI: 10.1016/j.rse.2023.113924. (Visited on 10/11/2024).

References VII

- [21] Jason Yosinski et al. *How Transferable Are Features in Deep Neural Networks?* Nov. 2014. DOI: 10.48550/arXiv.1411.1792. arXiv: 1411.1792 [cs]. (Visited on 01/20/2025).
- [22] Peng Liang, Wenzhong Shi, and Xiaokang Zhang. “Remote Sensing Image Classification Based on Stacked Denoising Autoencoder”. In: *Remote Sensing* 10.1 (Jan. 2018), p. 16. ISSN: 2072-4292. DOI: 10.3390/rs10010016. (Visited on 01/20/2025).
- [23] Lovedeep Gondara. “Medical Image Denoising Using Convolutional Denoising Autoencoders”. In: *2016 IEEE 16th International Conference on Data Mining Workshops (ICDMW)*. Dec. 2016, pp. 241–246. DOI: 10.1109/ICDMW.2016.0041. (Visited on 01/20/2025).

- [24] Pascal Vincent et al. “Extracting and Composing Robust Features with Denoising Autoencoders”. In: *Proceedings of the 25th International Conference on Machine Learning*. ICML '08. New York, NY, USA: Association for Computing Machinery, July 2008, pp. 1096–1103. ISBN: 978-1-60558-205-4. DOI: 10.1145/1390156.1390294. (Visited on 01/19/2025).